

# Cours X : understanding code performance

Charles Bouillaguet  
(`charles.bouillaguet@lip6.fr`)

2022-XX-XXXX

# Premature optimization is the root of all evil — Donald Knuth

## Meaning

- ▶ Optimization takes (programmer) time
- ▶ Use your time efficiently
- ▶ Optimization usually makes code more complex
  - ▶ Harder to read/debug/maintain/upgrade
- ▶ Should be done **only** when necessary

## Rule of thumb

- ▶ 95% of the running time...
  - ▶ ... usually concentrated in 5% of code statements
    - ▶ E.g. inner loops, specific function, etc.
- ~> Need to identify them

## Need

- ▶ We need tools to **observe** code behavior
- ▶ **Where** does it spend most of its time?

## Key idea

- ▶ **do not** optimize code representing 0.001% of the time
- ▶ Common rookie mistake

# Where Does my Code Spend Its Time?

## Two methods

1. Sampling
2. Instrumentation

## Illustration / Case Study

- ▶ Colleague from the *Institut de minéralogie, de physique des matériaux et de cosmochimie* (Sorbonne University)
- ▶ Least square fitting of model to experimental data points
- ▶ **Function**  $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ 
  - ▶  $f$  is computed by X million lines of ugly Fortran code
  - ▶  $n = 1000$  (parameters to choose)
  - ▶  $m = 10^7$  (data points to fit)
- ▶ **Goal:** find  $x$  that minimizes  $\|f(x)\|_2$
- ▶ Use MINPACK library from 1978 — **takes forever**. Why?

# Sampling

## Main Idea

- ▶ Just **run** the code
- ▶ **Interrupt** it regularly
  - ▶ E.g. 1kHz
- ▶ Store the **instruction pointer** ("sample")
- ▶ Resume

## Afterwards

- ▶ Look **where** it was interrupted
- ▶ Instruction often executed  $\rightsquigarrow$  more samples

Samples	92ca:	mov	%ecx,%edx
	92cc:	sub	0x8(%rsp),%r9
	92d1:	xor	%eax,%eax
	92d3:	shr	%edx
	92d5:	shl	\$0x4,%rdx
2	92d9:	nopl	0x0(%rax)
3	92e0:	movupd	(%r14,%rax,1),%xmm2
525	92e6:	movupd	(%r9,%rax,1),%xmm5
924	92ec:	add	\$0x10,%rax
28	92f0:	mulpd	%xmm5,%xmm2
2932	92f4:	addsd	%xmm2,%xmm0
2	92f8:	unpckhpd	%xmm2,%xmm2
3952	92fc:	addsd	%xmm2,%xmm0
	9300:	cmp	%rax,%rdx
4	9303:	jne	92e0 <qrfac_+0x3b0>
	9305:	mov	0x1c(%rsp),%eax
	9309:	mov	%ecx,%edx
2	930b:	and	\$0xffffffff,%edx
	930e:	add	%edx,%eax
	9310:	cmp	%edx,%ecx
	9312:	jne	932e <qrfac_+0x3fe>
	9314:	lea	0x0(%rbp,%rax,1),%edx
	9318:	add	%r12d,%eax
	931b:	movslq	%edx,%rdx
	931e:	cltq	
	9320:	movsd	(%rbx,%rdx,8),%xmm1
	9325:	mulsd	(%rbx,%rax,8),%xmm1
3	932a:	addsd	%xmm1,%xmm0
	932e:	mov	0x20(%rsp),%rax
1	9333:	lea	(%rsi,%r13,1),%r9
17	9337:	divsd	(%rax),%xmm0
	933b:	lea	0x0(,%r9,8),%rax

# Case Study

Percent	92d3:	shr	%edx
	92d5:	shl	\$0x4,%rdx
0,02	92d9:	nopl	0x0(%rax)
0,03	92e0:	movupd	(%r14,%rax,1),%xmm2
4,87	92e6:	movupd	(%r9,%rax,1),%xmm5
8,70	92ec:	add	\$0x10,%rax
0,26	92f0:	mulpd	%xmm5,%xmm2
27,60	92f4:	addsd	%xmm2,%xmm0
0,02	92f8:	unpckhpd	%xmm2,%xmm2
37,24	92fc:	addsd	%xmm2,%xmm0
	9300:	cmp	%rax,%rdx
0,04	9303:	jne	92e0 <qrfac +0x3b0>
	9305:	mov	0x1c(%rsp),%eax

## Pinpointing the Hot Spot

- ▶ Loop of 9 CPU instructions
- ▶  $\approx 80\%$  of the total running time

## Good, but...

- ▶ Where is instruction 0x92fc in my code? hint: qrfac

# Making Sense of Code Addresses

Where is instruction 0x92fc in my code?

## Good News

- ▶ The **debugging symbols** exist to answer this question
- ▶ Enable debugging symbols: `gcc [...] -g [...]`
  - ▶ Makes executable slightly larger...
- ▶ Map between **instructions** and **source file/line**
  - ▶ Sometimes not obvious with compiler optimizations
  - ▶ May reorder code, permute loops, etc.
- ▶ Exploited by many tools: `gdb`, `valgrind`, ...

## Simple tool: `addr2line`

```
$ addr2line -e ./speed_lmdif1 0x92fc  
minpack/qrfac.c:134 (discriminator 3)
```

Let's look at `minpack/qrfac.c`, line 134

## Actually looking at the code

```
127 void qrfac_(...)
128 {
129     for (int j = 1; j <= n; ++j) {
130         // ...
131         for (int k = j + 1; k <= n; ++k) {
132             double sum = 0;
133             for (int i = j; i <= m; ++i)
134 /* !!! */          sum += a[i + j * m] * a[i + k * m];
135             double temp = sum / a[j + j * m];
136             for (int i = j; i <= m; ++i)
137                 a[i + k * m] -= temp * a[i + j * m];
138             // ...
139         }
140         // ...
141     }
142 }
```



# Sampling — Summary

## (Relative) Ease of Use

- ▶ Commercial programs (Intel VTune, ...)
- ▶ Under linux: perf
  - ▶ `$ perf record ./program`
  - ▶ `$ perf report`

## Advantages

- ▶ Runs at  $\approx 100\%$  native speed
- ▶ Quite precise
- ▶ Can measure other things than time

## Problems

- ▶ Usually requires **administrator** privileges
- ▶ And/or cooperation from the operating system

# Instrumentation

## Main Idea

- ▶ **Modify** the code
  - ▶ Add measurements
- ▶ Run **instrumented** code

```
double start = wtime();
double qrfac_time = 0;
int qrfac_calls = 0;
// ...
double qrfac_start = wtime();
qrfac(...);
qrfac_time += wtime() - qrfac_start;
qrfac_calls += 1;
// ...
double total = wtime() - start;
printf("Running time: %.1fs\n", total);
printf("%d calls to qrfac (%.1fs)\n",
       qrfac_calls, qrfac_time);
exit(0);
```

## Afterwards

- ▶ Look at the data collected by the instrumentation

## [Manual | Automatic] instrumentation

- ▶ Manual: if you know what you're looking for
  - ▶ Or if you want your program to print performance results
- ▶ Automatic: more heavyweight, simpler

# Instrumentation: Goals

## What data do we obtain in the end?

### ▶ A **Profile**

- ▶ Short summary of performance results
- ▶ Typically per-function
- ▶ Eventually call-graph informations
  - ▶ "Function A takes 95% of the time, ..."
  - ▶ "...but only when it is called by function B"

### ▶ A **Trace**

- ▶ Log of timestamped "events"
- ▶ Visual view of performance problems
- ▶ Very powerful, more complex to use

Can build a profile from a trace, but not the other way around

# Instrumentation: the GNU Profiler

Easy and Always Available

## Usage

- ▶ Compile **and** link with: `gcc [...] -pg [...]`
    - ▶ Automatic instrumentation
  - ▶ Run program
- ~> `gmon.out`
- ▶ `$ gprof ./program`

# GNU Profiler: Result

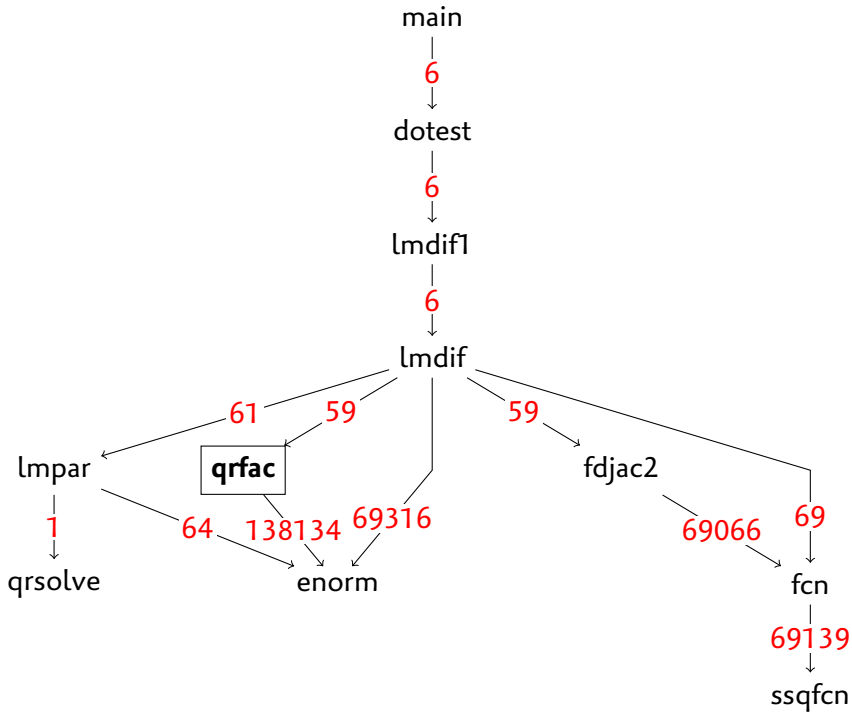
## Flat profile

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
97.91	408.82	408.82	59	6.93	6.95	qrfac_
0.91	412.63	3.81	1	3.81	3.81	qrsolv_
0.48	414.64	2.01	69139	0.00	0.00	ssqfcn
0.41	416.35	1.71	207520	0.00	0.00	enorm_
0.17	417.05	0.70	6	0.12	69.59	lmdif_
0.11	417.49	0.44	59	0.01	0.04	fdjac2_
0.02	417.56	0.07	61	0.00	0.06	lmpar_
0.00	417.56	0.00	69133	0.00	0.00	fcn
0.00	417.56	0.00	14	0.00	0.00	wtime
0.00	417.56	0.00	6	0.00	69.59	do_test
0.00	417.56	0.00	6	0.00	0.00	initpt
0.00	417.56	0.00	6	0.00	69.59	lmdif1_

- ▶ We learn that 98% of the time is spent in qrfac
- ▶ Called 59 times,  $\approx 6.9$ s per call

# Call Graph

	index	% time	self	children	called	name
	...					
			0.70	416.86	6/6	lmdif1_ [3]
[4]	100.0	0.70	416.86	6	lmdif_ [4]	
		408.82	1.14	59/59	qrfac_ [5]	
		0.07	3.81	61/61	lmpar_ [6]	
		0.44	2.01	59/59	fdjac2_ [8]	
		0.57	0.00	69316/207520	enorm_ [11]	
		0.00	0.00	67/69133	fc_ [10]	
		408.82	1.14	59/59	lmdif_ [4]	
[5]	98.2	408.82	1.14	59	qrfac_ [5]	
		1.14	0.00	138134/207520	enorm_ [11]	
		0.07	3.81	61/61	lmdif_ [4]	
[6]	0.9	0.07	3.81	61	lmpar_ [6]	
		3.81	0.00	1/1	qrsolv_ [7]	
		0.00	0.00	64/207520	enorm_ [11]	
		3.81	0.00	1/1	lmpar_ [6]	
[7]	0.9	3.81	0.00	1	qrsolv_ [7]	
		0.44	2.01	59/59	lmdif_ [4]	
[8]	0.6	0.44	2.01	59	fdjac2_ [8]	
		0.00	2.01	69066/69133	fc_ [10]	
	...					



# Profiling: Summary

## Not just gprof

- ▶ score-p, scalasca, tau, ...
- ▶ Profile more than time (bytes send/s, cache miss, ...)

## Advantages

- ▶ Ease of use (just recompile), **Good first approach**
- ▶ No need for root privileges
- ▶ Call-graph information

## Problems

- ▶ Coarse grained
  - ▶ Does not precisely pinpoint a specific loop
- ▶ May slow program down



# Traces



## Multi-thread

- ▶ Load imbalance (inactive threads)
- ▶ Time spent waiting for sync. (locks, barriers, ...)

# Traces



## MPI

- ▶ Time spent in communications
- ▶ What process delays the others?

# Traces

## Advantages

- ▶ Very precise information about dynamic code behavior
- ▶ Can be obtained automatically (e.g. score-p instrumenter)

## Problems

- ▶ Visualization tools are mostly commercial (Vampir, itac)
- ▶ Traces can be **large** with many processes/threads